

Machine Learning Algebra

By Matthieu Lagarde

April 9, 2022

1 Multivariate linear regression

1.1 Notations

Let m be the number of examples or observations in our training set. Let n be the number of features or explanatory variables observed for each example of the training set. Let $x_j^{(i)}$ be the value of feature j for example i . Let $y^{(i)}$ be the output value for example i .

$$y = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(m)} \end{pmatrix} \in \mathbb{R}^m \quad (1)$$

y is called the output vector. It contains the output values of the training set.

$$X = \begin{pmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_n^{(2)} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_1^{(m)} & x_2^{(m)} & \dots & x_n^{(m)} \end{pmatrix} \in \mathbb{R}^{m \times (n+1)} \quad (2)$$

X is called the data matrix. If we ignore the first column of 1, each row of matrix X is one example of the training set and each column of the matrix X is the values observed for one feature in the training set.

$$\theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \dots \\ \theta_n \end{pmatrix} \in \mathbb{R}^{(n+1)} \quad (3)$$

θ is the vector of parameters.

1.2 Hypothesis

We assume that there is a linear relationship between the features and the output. Note that the relationship is linear in θ but the features themselves can be non linear transformations of the initial features such as quadratic terms or interaction terms.

The unvectorized form of the hypothesis for a given example i is:

$$h_{\theta}(x^{(i)}) = \sum_{j=0}^n \theta_j * x_j^{(i)} \in \mathbb{R} \text{ with } x_0^{(i)} = 1 \quad (4)$$

The vectorized form of the hypothesis can be written as follows:

$$h_{\theta}(X) = X\theta \in \mathbb{R}^m \quad (5)$$

For a single example, we can also write the vectorized form of the hypothesis:

$$x^{(i)} = \begin{pmatrix} 1 \\ x_1^{(i)} \\ \dots \\ x_n^{(i)} \end{pmatrix} \in \mathbb{R}^{(n+1)}$$

$$h_{\theta}(x^{(i)}) = x^{(i)\top} \theta \in \mathbb{R}$$

1.3 Cost function

The unvectorized form of the cost function is:

$$J(\theta) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \in \mathbb{R}$$

where $x^{(i)} \in \mathbb{R}^{(n+1)}$ and

where $y^{(i)} \in \mathbb{R}$

The vectorized form of the cost function is:

$$J(\theta) = \frac{1}{2m} * (X\theta - y)^{\top} (X\theta - y) \in \mathbb{R} \quad (6)$$

1.4 Normal equation

$J(\theta)$ is convex so any local minimum is a global minimum. We thus know that θ^* obeys the following equation:

$$\nabla J(\theta^*) = 0_{(n+1)}$$

Let recall a property of matrix differentiation. Let α be a scalar equal to $y^\top x$ where y and x be two column vectors of \mathbb{R}^m that are respectively a function of another column vector z of \mathbb{R}^n . We have:

$$\alpha = y^\top x$$

$$\frac{\partial \alpha}{\partial z} = \frac{\partial y^\top}{\partial z} x + \frac{\partial x^\top}{\partial z} y \in \mathbb{R}^n$$

Using this property, we have:

$$\begin{aligned} 1/2m * 2 * \frac{\partial (X\theta^* - y)^\top}{\partial \theta} (X\theta^* - y) &= 0_{(n+1)} \\ \iff 1/m * X^\top (X\theta^* - y) &= 0_{(n+1)} \\ \iff X^\top X\theta^* - X^\top y &= 0_{(n+1)} \\ \boxed{\iff \theta^* = (X^\top X)^{-1} X^\top y \in \mathbb{R}^{(n+1)}} \end{aligned}$$

1.5 Gradient descent

If necessary, here is the vectorized implementation of gradient descent. Denoting α the learning rate:

$$\theta := \theta - \frac{\alpha}{m} * X^\top (X\theta - y) \in \mathbb{R}^{(n+1)} \quad (7)$$

1.6 Regularized cost function

Denoting λ the regularization parameter, the unvectorized form of the cost function can be written as:

$$J(\theta) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} * \sum_{j=1}^n \theta_j^2 \in \mathbb{R}$$

where $x^{(i)} \in \mathbb{R}^{(n+1)}$ and

where $y^{(i)} \in \mathbb{R}$

Be careful, in the regularization part of the expression (the second sum), the j index goes from 1 to n and NOT from 0 to n . Indeed, by convention, we do not regularize θ_0 .

The vectorized form of the regularized cost function is thus:

$$J(\theta) = \frac{1}{2m} * (X\theta - y)^{\top} (X\theta - y) + \frac{\lambda}{2m} * \theta_r^{\top} \theta_r$$

$$\text{where } \theta_r = \begin{pmatrix} \theta_1 \\ \dots \\ \theta_n \end{pmatrix} \in \mathbb{R}^n$$

1.7 Regularized normal equation

The vectorized form of the regularized normal equation is:

$$\theta^* = (X^{\top} X + \lambda * M)^{-1} X^{\top} y \in \mathbb{R}^{(n+1)}$$

$$\text{where } M = \begin{pmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & 1 & 0 \\ 0 & 0 & \dots & 0 & 1 \end{pmatrix} \in \mathbb{R}^{(n+1) \times (n+1)}$$

1.8 Regularized gradient descent

Let compute the gradient of $J(\theta)$ when it is regularized. There are two cases, one for the partial derivative with respect to θ_0 and one for the partial derivatives with respect to θ_j :

$$\frac{\partial J(\theta)}{\partial \theta_0} = \left[\frac{1}{m} * X^\top (X\theta - y) \right]_1 \text{ i.e. 1st element of previous gradient}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \left[\frac{1}{m} * X^\top (X\theta - y) + \frac{\lambda}{m} * \theta \right]_{j+1} \text{ for } j \in \{1, 2, \dots, n\}$$

2 Logistic regression

2.1 Sigmoid function

Let introduce the sigmoid function:

$$g : x \in \mathbb{R} \longrightarrow \frac{1}{1 + e^{-x}} \in]0, 1[\quad (8)$$

The interesting properties of the sigmoid function are:

- It is defined over \mathbb{R} .
- It is increasing.
- $g(0) = \frac{1}{2}$.
- It is converging to 0 in $-\infty$ and to 1 in $+\infty$.
- It is convex over $]-\infty, 0]$ and concave over $[0, +\infty[$.
- It means that (0, 0.5) is an inflection point of g .
- $g(-4) = 0.02$ and $g(4) = 0.98$




Figure 1: Graph of the sigmoid function

2.2 Hypothesis

The unvectorized form of the hypothesis for a given example i is:

$$h_{\theta}(x^{(i)}) = g\left(\sum_{j=0}^n \theta_j * x_j^{(i)}\right) \in]0, 1[$$

where g is the sigmoid function.

The hypothesis can be interpreted as the probability that a new observation belongs to the positive class i.e. that $y = 1$ given the value of its features i.e. given x , parameterized by θ . In other words:

$$h_{\theta}(x) = P(y = 1|x; \theta) \tag{9}$$

We then introduce the following decision rule:

$$y = \begin{cases} 1 & \text{if } h_{\theta}(x) \geq 0.5 \\ 0 & \text{if } h_{\theta}(x) < 0.5 \end{cases} \tag{10}$$

Finally, the vectorized form of the hypothesis can be written as follows:

$$h_{\theta}(X) = g(X\theta) \in]0, 1[^m$$

where g is the sigmoid function applied element-wise

2.3 Cost function

Before introducing the cost function, let study two functions:

$$v : x \in]0, 1[\longrightarrow -\ln(x) \in [0, +\infty[$$

$$w : x \in]0, 1[\longrightarrow -\ln(1 - x) \in [0, +\infty[$$

Function v has the following properties on the interval $]0, 1[$:

- It is decreasing.
- It converges to $+\infty$ in 0.
- $v(1) = 0$.

Function w has the following properties on the interval $]0, 1[$:

- It is increasing.
- It converges to $+\infty$ in 1.
- $w(0) = 0$.

Here is the graph of v and w on the interval $]0, 1[$:




Figure 2: Graph of v and w

Now we introduce the following cost function:

$$J(\theta) = \frac{1}{m} * \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \in \mathbb{R}$$

where

$$\text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) = \begin{cases} -\ln(h_{\theta}(x^{(i)})) & \text{if } y^{(i)} = 1 \\ -\ln(1 - h_{\theta}(x^{(i)})) & \text{if } y^{(i)} = 0 \end{cases}$$

Given that $y \in \{0, 1\}$, the cost function can be rewritten as:

$$J(\theta) = \frac{1}{m} * \sum_{i=1}^m [(y^{(i)} * -\ln(h_{\theta}(x^{(i)}))) + ((1 - y^{(i)}) * -\ln(1 - h_{\theta}(x^{(i)})))] \in \mathbb{R} \quad (11)$$

We can also write the vectorized form of the cost function as follows:

$$J(\theta) = -\frac{1}{m} * [y^{\top} \ln(h_{\theta}(X)) + (1 - y)^{\top} \ln(1 - h_{\theta}(X))] \in \mathbb{R}$$

where \ln is applied element-wise

2.4 Gradient descent

The unvectorized form of the gradient of the cost function can be written as:

$$\frac{\partial J(\theta)}{\partial \theta_j} = \quad (12)$$