

Machine Learning Algebra

By Matthieu Lagarde

April 9, 2022

1 Multivariate linear regression

1.1 Notations

Let m be the number of examples or observations in our training set. Let n be the number of features or explanatory variables observed for each example of the training set. Let $x_j^{(i)}$ be the value of feature j for example i . Let $y^{(i)}$ be the output value for example i .

$$y = \begin{pmatrix} y^{(1)} \\ y^{(2)} \\ \dots \\ y^{(m)} \end{pmatrix} \in \mathbb{R}^m \quad (1)$$

y is called the output vector. It contains the output values of the training set.

$$X = \begin{pmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \dots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \dots & x_n^{(2)} \\ \dots & \dots & \dots & \dots & \dots \\ 1 & x_1^{(m)} & x_2^{(m)} & \dots & x_n^{(m)} \end{pmatrix} \in \mathbb{R}^{m \times (n+1)} \quad (2)$$

X is called the data matrix. If we ignore the first column of 1, each row of matrix X is one example of the training set and each column of the matrix X is the values observed for one feature in the training set.

$$\theta = \begin{pmatrix} \theta_0 \\ \theta_1 \\ \dots \\ \theta_n \end{pmatrix} \in \mathbb{R}^{(n+1)} \quad (3)$$

θ is the vector of parameters.

1.2 Hypothesis

We assume that there is a linear relationship between the features and the output. Note that the relationship is linear in θ but the features themselves can be non linear transformations of the initial features such as quadratic terms or interaction terms.

The unvectorized form of the hypothesis for a given example i is:

$$h_{\theta}(x^{(i)}) = \sum_{j=0}^n \theta_j * x_j^{(i)} \in \mathbb{R} \text{ with } x_0^{(i)} = 1 \quad (4)$$

The vectorized form of the hypothesis can be written as follows:

$$h_{\theta}(X) = X\theta \in \mathbb{R}^m \quad (5)$$

For a single example, we can also write the vectorized form of the hypothesis:

$$x^{(i)} = \begin{pmatrix} 1 \\ x_1^{(i)} \\ \dots \\ x_n^{(i)} \end{pmatrix} \in \mathbb{R}^{(n+1)}$$

$$h_{\theta}(x^{(i)}) = x^{(i)\top} \theta \in \mathbb{R}$$

1.3 Cost function

The unvectorized form of the cost function is:

$$J(\theta) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \in \mathbb{R}$$

where $x^{(i)} \in \mathbb{R}^{(n+1)}$ and

where $y^{(i)} \in \mathbb{R}$

The vectorized form of the cost function is:

$$J(\theta) = \frac{1}{2m} * (X\theta - y)^{\top} (X\theta - y) \in \mathbb{R} \quad (6)$$

1.4 Normal equation

$J(\theta)$ is convex so any local minimum is a global minimum. We thus know that θ^* obeys the following equation:

$$\nabla J(\theta^*) = 0_{(n+1)}$$

Let recall a property of matrix differentiation. Let α be a scalar equal to $y^\top x$ where y and x be two column vectors of \mathbb{R}^m that are respectively a function of another column vector z of \mathbb{R}^n . We have:

$$\alpha = y^\top x$$

$$\frac{\partial \alpha}{\partial z} = \frac{\partial y^\top}{\partial z} x + \frac{\partial x^\top}{\partial z} y \in \mathbb{R}^n$$

Using this property, we have:

$$\begin{aligned} 1/2m * 2 * \frac{\partial (X\theta^* - y)^\top}{\partial \theta} (X\theta^* - y) &= 0_{(n+1)} \\ \iff 1/m * X^\top (X\theta^* - y) &= 0_{(n+1)} \\ \iff X^\top X\theta^* - X^\top y &= 0_{(n+1)} \\ \boxed{\iff \theta^* = (X^\top X)^{-1} X^\top y \in \mathbb{R}^{(n+1)}} \end{aligned}$$

1.5 Gradient descent

If necessary, here is the vectorized implementation of gradient descent. Denoting α the learning rate:

$$\theta := \theta - \frac{\alpha}{m} * X^\top (X\theta - y) \in \mathbb{R}^{(n+1)} \quad (7)$$

1.6 Regularized cost function

Denoting λ the regularization parameter, the unvectorized form of the cost function can be written as:

$$J(\theta) = \frac{1}{2m} * \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \frac{\lambda}{2m} * \sum_{j=1}^n \theta_j^2 \in \mathbb{R}$$

where $x^{(i)} \in \mathbb{R}^{(n+1)}$ and

where $y^{(i)} \in \mathbb{R}$

Be careful, in the regularization part of the expression (the second sum), the j index goes from 1 to n and NOT from 0 to n . Indeed, by convention, we do not regularize θ_0 .

The vectorized form of the regularized cost function is thus:

$$J(\theta) = \frac{1}{2m} * (X\theta - y)^{\top} (X\theta - y) + \frac{\lambda}{2m} * \theta_r^{\top} \theta_r$$

$$\text{where } \theta_r = \begin{pmatrix} \theta_1 \\ \dots \\ \theta_n \end{pmatrix} \in \mathbb{R}^n$$

1.7 Regularized normal equation

The vectorized form of the regularized normal equation is:

$$\theta^* = (X^{\top} X + \lambda * M)^{-1} X^{\top} y \in \mathbb{R}^{(n+1)}$$

$$\text{where } M = \begin{pmatrix} 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & \dots & 1 & 0 \\ 0 & 0 & \dots & 0 & 1 \end{pmatrix} \in \mathbb{R}^{(n+1) \times (n+1)}$$

1.8 Regularized gradient descent

Let compute the gradient of $J(\theta)$ when it is regularized. There are two cases, one for the partial derivative with respect to θ_0 and one for the partial derivatives with respect to θ_j :

$$\frac{\partial J(\theta)}{\partial \theta_0} = \left[\frac{1}{m} * X^\top (X\theta - y) \right]_1 \text{ i.e. 1st element of previous gradient}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \left[\frac{1}{m} * X^\top (X\theta - y) + \frac{\lambda}{m} * \theta \right]_{j+1} \text{ for } j \in \{1, 2, \dots, n\}$$

2 Logistic regression

2.1 Sigmoid function

Let introduce the sigmoid function:

$$g : x \in \mathbb{R} \longrightarrow \frac{1}{1 + e^{-x}} \in]0, 1[\quad (8)$$

The interesting properties of the sigmoid function are:

- It is defined over \mathbb{R} .
- It is increasing.
- $g(0) = \frac{1}{2}$.
- It is converging to 0 in $-\infty$ and to 1 in $+\infty$.
- It is convex over $]-\infty, 0]$ and concave over $[0, +\infty[$.
- It means that (0, 0.5) is an inflection point of g .
- $g(-4) = 0.02$ and $g(4) = 0.98$




Figure 1: Graph of the sigmoid function

2.2 Hypothesis

The unvectorized form of the hypothesis for a given example i is:

$$h_{\theta}(x^{(i)}) = g\left(\sum_{j=0}^n \theta_j * x_j^{(i)}\right) \in]0, 1[$$

where g is the sigmoid function.

The hypothesis can be interpreted as the probability that a new observation belongs to the positive class i.e. that $y = 1$ given the value of its features i.e. given x , parameterized by θ . In other words:

$$h_{\theta}(x) = P(y = 1|x; \theta) \tag{9}$$

We then introduce the following decision rule:

$$y = \begin{cases} 1 & \text{if } h_{\theta}(x) \geq 0.5 \\ 0 & \text{if } h_{\theta}(x) < 0.5 \end{cases} \tag{10}$$

Finally, the vectorized form of the hypothesis can be written as follows:

$$h_{\theta}(X) = g(X\theta) \in]0, 1[^m$$

where g is the sigmoid function applied element-wise

2.3 Cost function

Before introducing the cost function, let study two functions:

$$v : x \in]0, 1[\longrightarrow -\ln(x) \in [0, +\infty[$$

$$w : x \in]0, 1[\longrightarrow -\ln(1 - x) \in [0, +\infty[$$

Function v has the following properties on the interval $]0, 1[$:

- It is decreasing.
- It converges to $+\infty$ in 0.
- $v(1) = 0$.

Function w has the following properties on the interval $]0, 1[$:

- It is increasing.
- It converges to $+\infty$ in 1.
- $w(0) = 0$.

Here is the graph of v and w on the interval $]0, 1[$:

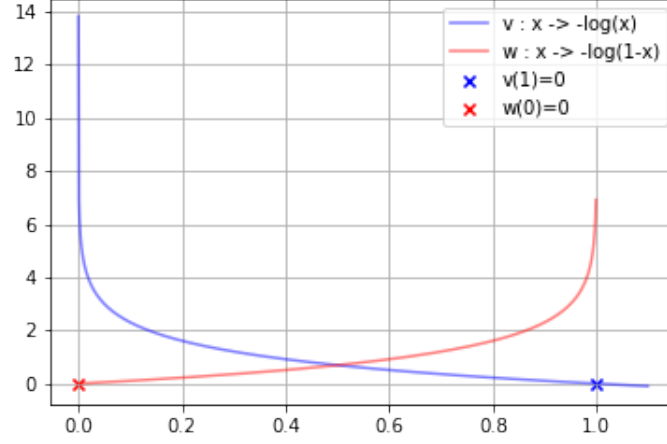


Figure 2: Graph of v and w

Now we introduce the following cost function:

$$J(\theta) = \frac{1}{m} * \sum_{i=1}^m Cost(h_{\theta}(x^{(i)}), y^{(i)}) \in \mathbb{R}$$

where

$$Cost(h_{\theta}(x^{(i)}), y^{(i)}) = \begin{cases} -\ln(h_{\theta}(x^{(i)})) & \text{if } y^{(i)} = 1 \\ -\ln(1 - h_{\theta}(x^{(i)})) & \text{if } y^{(i)} = 0 \end{cases}$$

Given that $y \in \{0, 1\}$, the cost function can be rewritten as:

$$J(\theta) = \frac{1}{m} * \sum_{i=1}^m [(y^{(i)} * -\ln(h_{\theta}(x^{(i)}))) + ((1 - y^{(i)}) * -\ln(1 - h_{\theta}(x^{(i)})))] \in \mathbb{R} \quad (11)$$

We can also write the vectorized form of the cost function as follows:

$$J(\theta) = -\frac{1}{m} * [y^{\top} \ln(h_{\theta}(X)) + (1 - y)^{\top} \ln(1 - h_{\theta}(X))] \in \mathbb{R}$$

where \ln is applied element-wise

2.4 Gradient descent

The unvectorized form of the gradient of the cost function can be written as:

$$\frac{\partial J(\theta)}{\partial \theta_k} = \frac{1}{m} * \sum_{i=1}^m \left(\frac{1}{1 + \exp(-\sum_{j=0}^n x_j^{(i)} * \theta_j)} - y^{(i)} \right) * x_k^{(i)} \in \mathbb{R} \quad (12)$$

The vectorized form of the gradient of the cost function can be written as:

$$\nabla J(\theta) = \frac{1}{m} * X^\top (h_\theta(X) - y) \in \mathbb{R}^{n+1}$$

where $h_\theta(X) \in]0, 1[^m$

Denoting α the learning rate, a vectorized implementation of the gradient descent can thus be written as:

$$\theta := \theta - \frac{\alpha}{m} * X^\top (h_\theta(X) - y) \in \mathbb{R}^{n+1}$$

where $h_\theta(X) = g(X\theta) \in]0, 1[^m$

and g is the sigmoid function applied element-wise

2.5 Regularization

The regularized cost function can be written as:

$$J(\theta) = -\frac{1}{m} * [y^\top \ln(h_\theta(X)) + (1 - y)^\top \ln(1 - h_\theta(X))] + \frac{\lambda}{2m} * \theta_r^\top \theta_r \in \mathbb{R}$$

where \ln is applied element-wise

The regularized gradient can be written as:

$$\begin{aligned} \left[\widetilde{\nabla J(\theta)} \right]_1 &= [\nabla J(\theta)]_1 \text{ for } \partial \theta_0 \text{ i.e. the first element of the regularized gradient} \\ \left[\widetilde{\nabla J(\theta)} \right]_j &= \left[\nabla J(\theta) + \frac{\lambda}{m} * \theta \right]_j \text{ where } j \in \{2, 3, \dots, n\} \end{aligned}$$

3 Neural network classifier

3.1 Notations

Let consider the following neural network:

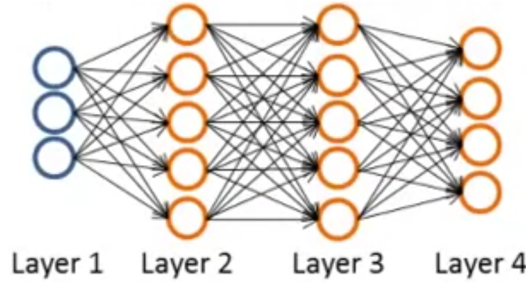


Figure 3: Example of neural network

We assume that we face a multi-class classification where we want to predict whether an observation belongs to one of K possible classes. We denote L the total number of layers in the network. In the example, there are 4 layers. The first layer is called the input layer. The last layer is called the output layer. Intermediate layers are called hidden layers. Here, there are 2 hidden layers. We denote s_l the number of units in layer l , not counting the bias unit. In the example, we thus have:

- $s_1 = 3$ which is n , the number of input features.
- $s_2 = 5$.
- $s_3 = 5$.
- $s_4 = 4$ which is K , the number of classes.

We assume that a bias unit is added as an input to any layer.

We denote m the total number of examples in our training set. We denote $x^{(i)} \in \mathbb{R}^{(s_1+1)}$ the vector of features of the i -th example in the training set where $x_0^{(i)} = 1$. We denote $y^{(i)} \in \mathbb{R}^{K=s_L}$ the output vector of the i -th example in the training set. If the i -th example belongs to class k , then:

$$y^{(i)} = \begin{pmatrix} 0 \\ \dots \\ 0 \\ 1 \\ 0 \\ \dots \\ 0 \end{pmatrix} \text{ where } (y^{(i)})_{k,1} = 1 \in \mathbb{R}^{K=s_L} \quad (13)$$

Note that the full output values can be stored in a matrix:

$$Y = \begin{pmatrix} y^{(1)\top} \\ \dots \\ y^{(m)\top} \end{pmatrix} \in \mathbb{R}^{m \times K} \quad (14)$$

Finally, we denote $\Theta^{(l)}$ the matrix of parameters required to go from layer l to layer $(l + 1)$. This matrix is in the space $\mathbb{R}^{s_{(l+1)} \times ((s_l)+1)}$ since it maps vectors of dimension $((s_l) + 1)$ (do not forget the bias unit) into vectors of dimension $s_{(l+1)}$. Note that:

$$\Theta^{(l)} = \begin{pmatrix} \text{Parameters associated to unit 1 of layer } (l + 1) \\ \dots \\ \text{Parameters associated to unit } s_{l+1} \text{ of layer } (l + 1) \end{pmatrix} \in \mathbb{R}^{s_{(l+1)} \times ((s_l)+1)} \quad (15)$$

In other words, $(\Theta^l)_{i,j}$ corresponds to the parameter associated to unit j of the layer l , which is used to compute unit i of the layer $(l + 1)$. Note that in this example, there are 3 matrices of parameters to estimate : (Θ^1) , (Θ^2) and (Θ^3) . It implies $5 * 3 + 5 * 5 + 4 * 5 = 15 + 25 + 20 = 60$ parameters to estimate. In general, there are $(L - 1)$ matrices of parameters to estimate.

3.2 Hypothesis

The hypothesis can be written recursively as follows:

$X_1 = \text{the data matrix} \in \mathbb{R}^{m \times (s_1+1)}$

$X_{l+1} = 1_{m \times 1} \parallel g(X_l \Theta^{(l)\top}) \in \mathbb{R}^{m \times ((s_{l+1})+1)}$

$h_{\Theta}(X_1) = g(X_{L-1} \Theta^{(L-1)\top}) \in \mathbb{R}^{m \times s_L}$

where \parallel means horizontal concatenation and g is the sigmoid function applied element-wise.

Again, the hypothesis only gives the probability that an observation belongs to each class. As for the one vs. all multi-class classification problem, we then introduce the following decision rule to pick the most probable class given the input features:

$$\text{Class of } y = \text{index}_{of} (\max_{row\ wise} (h_{\Theta}(X_1))) \in \mathbb{R}^m \quad (16)$$

3.3 Cost function

Very similar to the cost function for the logistic regression, we just add a cost for the estimation of each output unit (now there are K of them). We also directly introduce the regularization term. Here is an unvectorized expression of the cost function:

$$\begin{aligned} J(\Theta) = & -\frac{1}{m} * \sum_{i=1}^m \sum_{k=1}^K y_k^{(i)} * \ln(h_{\Theta}(X_1)_{i,k}) + (1 - y_k^{(i)}) * \ln(1 - h_{\Theta}(X_1)_{i,k}) \\ & + \frac{\lambda}{2m} * \sum_{l=1}^{L-1} \sum_{i=1}^{s_{(l+1)}} \sum_{j=2}^{(s_l)+1} \Theta_{i,j}^{(l)2} \end{aligned} \quad (17)$$

It can be vectorized as follows:

$$J(\Theta) = -\frac{1}{m} * [Tr (ln(h_{\Theta}(X_1))Y^{\top}) + Tr (ln(1 - h_{\Theta}(X_1))(1 - Y)^{\top})] \\ + \frac{\lambda}{2m} * \sum_{l=1}^{L-1} Tr (\Theta_r^{(l)\top} \Theta_r^{(l)})$$

where $\Theta_r^{(l)} = \Theta^{(l)}$ without the first column of θ_0 s,

and $h_{\Theta}(X_1) \in \mathbb{R}^{m \times K}$ and $Y \in \mathbb{R}^{m \times K}$,

and $Tr(.)$ is the trace of a square matrix operator i.e. the sum of its diagonal, $\ln(.)$ is applied element-wise.

3.4 Backpropagation algorithm

In order to minimize the cost function, we need to compute the partial derivatives of $J(\Theta)$ with respect to all the parameters i.e. with respect to each element of the $(L - 1)$ matrices of parameters:

$$\frac{\partial J(\Theta)}{\partial \Theta_{i,j}^{(l)}} \text{ for } i, j, l \quad (18)$$

To do so, we can use the backpropagation algorithm. This algorithm will yields the partial derivatives of $J(\Theta)$ with respect to each element of the matrices of parameters.

- Step 1: Initialization. Set the matrices of parameters to random small values.
- Step 2: Compute forward the activation of the units of each intermediate layer as well as the hypothesis i.e. the output units with matrices of zeroes.
- Step 3: Compute the error of the last layer.
- Step 4: Compute backward the errors of all previous layers until the first intermediate layer.
- Step 5: Use the errors of each layer to compute the gradients of each matrix of parameters.

On the full sample, the errors can be computed as follows:

$$\begin{aligned}
\delta^{(L)} &= A^{(L)} - Y^{(L)} \in \mathbb{R}^{m \times s_L} \\
\delta^{(l)} &= \delta^{(l+1)} \Theta^{(l)} \cdot * A^{(l)} \cdot * (1 - A^{(l)}) \in \mathbb{R}^{m \times (s_l+1)} \\
\text{where } A^{(l)} &= 1_{m \times 1} \parallel g(\Theta^{(l-1)} X_{l-1}^\top)
\end{aligned}$$

Finally, we can obtain the gradients as follows:

$$\begin{aligned}
\Delta^{(l)} &= \frac{1}{m} * \delta^{(l+1)\top} A^{(l)} + \frac{\lambda}{m} * \Theta_r^{(l)} \\
\text{where } \Theta_r^{(l)} &= \begin{pmatrix} 0 \\ 0 \\ \dots \\ 0 \end{pmatrix} \parallel [\Theta^{(l)}]_{col2-end}
\end{aligned}$$

Note that for $l = 1$, we have $A^1 = X$.

4 Support Vector Machines

4.1 Hypothesis

Support Vector Machines or SVM is another classifier. The hypothesis is a discriminant function based on the value of the linear combination:

$$\begin{aligned}
h_\theta(x^{(i)}) &= \begin{cases} 1 & \text{if } \theta^\top x^{(i)} \geq 0 \\ 0 & \text{otherwise} \end{cases} \\
\text{where } x^{(i)} &\in \mathbb{R}^{(n+1)} \text{ and } \theta \in \mathbb{R}^{(n+1)}
\end{aligned}$$

4.2 Cost function

The cost function is inspired by the cost function of the logistic regression but is slightly different:

$$J(\theta) = C * \left[\sum_{i=1}^m y^{(i)} * cost_1(\theta^\top x^{(i)}) + (1 - y^{(i)}) * cost_0(\theta^\top x^{(i)}) \right] + \frac{1}{2} \sum_{j=1}^n \theta_j^2 \quad (19)$$

Compared to the cost function of the logistic regression:

- We have dropped the $\frac{1}{m}$ factors
- We have replaced the parameter λ by a parameter $C = \frac{1}{\lambda}$. The bigger C, the more we want to fit exactly the data. The lower C, the more we want the parameters to be small.
- The individual cost function assigning a cost to the prediction is the Hinge loss function.

It can be written in a vectorized form as follows:

$$J(\theta) = C * \left[(Y^\top \text{cost}_1(X\theta)) + ((1 - Y)^\top \text{cost}_0(X\theta)) \right] + \frac{1}{2} * \theta_r^\top \theta_r \in \mathbb{R}$$

where θ_r is equal to θ without θ_0 thus $\in \mathbb{R}^n$

$\text{cost}_0(\cdot)$ and $\text{cost}_1(\cdot)$ are applied element-wise.

4.3 Hinge loss functions

For $y^{(i)} = 1$, the hinge loss function is:

$$\text{cost}_1(z) = \begin{cases} 0 & \text{if } z \geq 1 \\ k(1 - z) & \text{if } z < 1 \end{cases} \text{ with } k \text{ an arbitrary slope parameter } > 0 \quad (20)$$

For $y^{(i)} = 0$, the hinge loss function is:

$$\text{cost}_0(z) = \begin{cases} 0 & \text{if } z \leq -1 \\ k(1 + z) & \text{if } z > -1 \end{cases} \text{ with } k \text{ an arbitrary slope parameter } > 0 \quad (21)$$

Here is a graph of cost_0 and cost_1 on $[-2, 2]$ with $k = 2$:




Figure 4: Graph of $cost_0$ and $cost_1$ with $k = 2$

4.4 Large margin intuition

When ...